




ner and pos when nothing is capitalized

stephen mayhew, tatiana tsygankova, dan roth
university of pennsylvania
{mayhew, ttasya, danroth}@seas.upenn.edu



Problem: NER and POS models aren't robust to casing issues.

Experiments: try 6 different solutions, report scores on cased (C) and uncased (U) data.

Key takeaway: data augmentation is the best option.

Casing Robustness is a REAL ISSUE

NOW HEAR THIS: NAVY ABANDONS ALL CAPS

Official Communications, Long Written Large, Can Use Mixed Case; No Shouting

Modern Tools Have TROUBLE***

her mother, susannah, lived near paris

David's Darkest Moments Came **PERSON** After Goliath

The Internet gives LIMITED HELP

answered Jul 15 '17 at 20:38



Christopher Manning
8,545 • 27 • 43

Approach 1: Caseless models. W will work much better on all lowerca

Approach 2: Use the truecaser. V into formally edited capitalization. Y

Experimental Setting: modify only the data.
Models are standard BiLSTM-CRF.

	Train Data	Test Data
1.	<Do nothing>	<Do nothing>
2.	Lowercase	Lowercase
3.	Augment with Lowercase	<Do nothing>
3.5	Sample with Lowercase	<Do nothing>
4.	<Do nothing>	Truecase
5.	Truecase	Truecase

NER (CoNLL 2003 English)

Exp.	Test (C)	Test (U)	Avg
1. Cased	92.45	34.46	63.46
2. Uncased	89.32	89.32	89.32
3. C+U	91.67	89.31	90.49
3.5. Half Mixed	91.68	89.05	90.37
4. Truecase Test	82.93	82.93	82.93
5. Truecase All	90.25	90.25	90.25

POS (PTB English)

Exp.	Test (C)	Test (U)	Avg
1. Cased	97.85	88.66	93.26
2. Uncased	97.45	97.45	97.45
3. C+U	97.79	97.35	97.57
3.5. Half Mixed	97.85	97.36	97.61
4. Truecase Test	95.21	95.21	95.21
5. Truecase All	97.38	97.38	97.38

Mention Detection (Twitter, English)

Exp.	Mention Detection F1
1. Cased	58.63
2. Uncased	53.13
3. C+U	66.14
3.5. Half Mixed	64.69
4. Truecase Test	58.22
5. Truecase All	62.66